

Quiz 1

● Graded

Student

Boning Li

Total Points

25 / 30 pts

Question 1

Question 1

6 / 6 pts

✓ - 0 pts Correct

Question 2

Question 2

4 / 6 pts

✓ - 2 pts Answer is correct (NO) but not firmly demonstrated by justification/counterexample

Question 3

Question 3

10 / 10 pts

✓ - 0 pts Correct

Question 4

Question 4

5 / 8 pts

✓ - 3 pts One answer not computed (algebra shown)

Distributed Systems and Algorithms — CSCI 4510/6510

Quiz 1

September 12, 2024

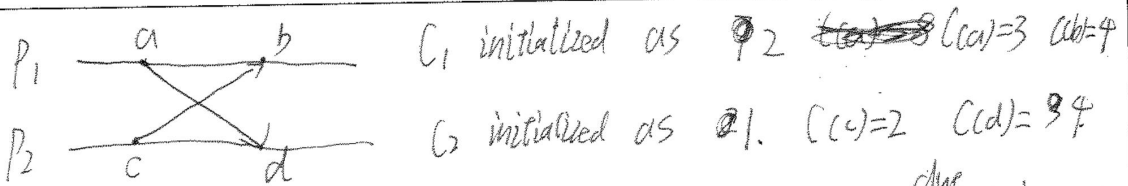
RCS ID: 62619 @rpi.edu Name: Boping Li

Instructions:

- You will have **45 minutes** to complete this quiz. Please do not start until told.
- Write your RCS ID and name in the blanks at the top of this cover sheet.
- Put away notes, laptops, and other electronic devices. Cheating on a quiz will result in an **immediate F** and a report will be filed with the Dean of Students.
- Read each question carefully several times before beginning to work and especially before asking questions.
- Write your answers clearly and completely inside the box.

Question 1 (6 points). Recall that in Lamport's Logical Clock algorithm, each process p_i has an integer c_i that is initialized to 0. Suppose instead, each process's clock is initialized to a random number between 0 where N is the number of processes in the system. The algorithm remains otherwise unchanged.

Does this modified algorithm always satisfy the weak clock condition, i.e., if $e \rightarrow f$ then $C(e) < C(f)$? Answer YES or NO and justify your answer.



If e and f occur on the same process P_i , then $C(e) < C(f)$ ~~due~~ to increment.

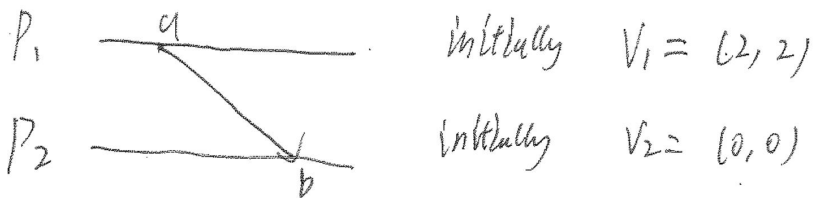
If e occurs at P_2 and f occurs at P_1 , then $C(f) = \max(C(e)+1, C_j(f))$,
 as a result, $C(e) < C(f)$.

Therefore, this modified algorithm always satisfies the weak clock condition.
 The answer is YES.

Question 2 (6 points). Recall that in Vector Clock algorithm, each process p_i has a vector clock VT_i that is initialized to the 0 vector. Suppose instead, each process's clock is initialized to be a vector of random numbers between 0 and N , where N is the number of processes in the system. The algorithm remains otherwise unchanged.

Does this modified algorithm always satisfy the strong clock condition, i.e., $e \rightarrow f$ if and only if $C(e) < C(f)$? Answer YES or NO and justify your answer.

In case where e and f occur on the same p_i , $e \rightarrow f$ gives $C(e) < C(f)$ since $e.VT(i) < f.VT(i)$ due to increment, and for all other $j \neq i$, $e.VT(j) \leq f.VT(j)$ due to inheritance. When e occurs on p_i and f on p_j , consider the following case.



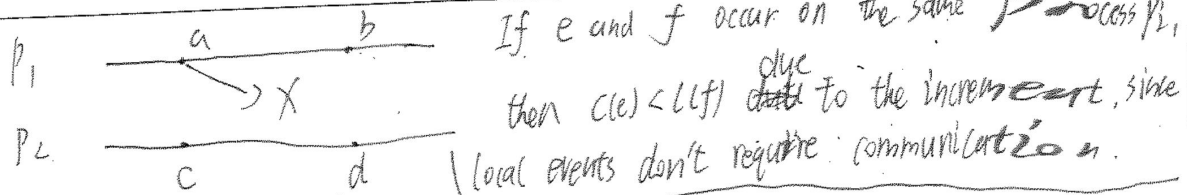
On event a , $V_1 = (3, 2)$ and is sent to P_2 .

Upon receiving in on event b , $b.VT = \max[(0, 1), (3, 2)] = (3, 2)$

Then we have $a \rightarrow b$ but $a.VT = b.VT$.

Therefore, the answer is NO.

Question 3 (10 points). All of the space-time diagrams I have drawn in class illustrate ~~reliable~~ ~~common~~ ~~communication~~ i.e., every message that is sent is eventually received. Consider a distributed system in which ~~common~~ ~~communication~~ ~~unreliable~~; some messages may be sent but never received (messages may be lost). Does Lamport's ~~Logi~~ ~~your~~ algorithm always satisfy the weak clock condition in this system: if $e \rightarrow f$ then $C(e) < C(f)$? Justify ~~your~~.



If e occurs on P_i and f occurs on P_j , then $e \rightarrow f$ states that e causally affects f . When e is a send of m , and f is the receive of m . In the case where f arrives at P_j , we have $C(e) < C(f)$ based on the algorithm. However, if m get lost, there won't be a receive event f on P_j , since P_j has no idea that P_i sends it a msg. In this case, there's no event f and no such relation $e \rightarrow f$ exists. As a result, as long as we have relation $e \rightarrow f$, $C(e) < C(f)$ will hold. Lamport's Logical clock algorithm still satisfies the weak condition.

Question 4 (8 points). Suppose the push algorithm is used to set the time at process p_i 's physical clock. The communication link between the time server S and process p_i is synchronous, with a minimum latency of 10 and a maximum latency of 30. The time server S sends its message to p_i at time $t = 100$. Answer the following questions and show your work.

1. When p_i receives the message, what does it set its clock to?
2. What is the resulting accuracy of p_i 's clock?

① Assume ~~the~~ the round-trip time recorded by p_i is $TRT \geq 20$.
 The ~~minimum~~ time ~~required~~ ^{earliest} for the message to arrive at p_i is 110.
 The latest time of arrival is $100 + 30 = 130$, or $TRT - \min = TRT - 10$
 A reasonable setup could be $\frac{130 + 110}{2} = 120$
 If $TRT - 10 < 30$, the interval becomes $[t + \min, t + TRT - \min] = [110, TRT + 100]$
 ② In this case, the reasonable setup should be $t + \frac{TRT}{2} = 100 + \frac{TRT}{2}$

② Given the time interval, the resulting accuracy is

$$\pm \left(\frac{TRT}{2} - \min \right) = \pm \left(\frac{TRT}{2} - 10 \right)$$

