

Quiz 2

● Graded

Student

Boning Li

Total Points

27 / 30 pts

Question 1

Question 1

10 / 10 pts

✓ - 0 pts Correct

Question 2

Question 2

10 / 10 pts

✓ - 0 pts Correct - solid justifications

Question 3

Question 3

7 / 10 pts

Correct response (NO)

✓ - 3 pts Correct response (NO), but justification is incorrect - execution does not follow algorithm

Distributed Systems and Algorithms — CSCI 4510/6510

Quiz 2

September 30, 2024


RCS ID: lib19 @rpi.edu Name: Boning Li

Instructions:

- You will have **45 minutes** to complete this quiz. Please do not start until told.
- Write your RCS ID and name in the blanks at the top of this cover sheet.
- Put away notes, laptops, and other electronic devices. Cheating on a quiz will result in an **immediate F** and a report will be filed with the Dean of Students.
- Read each question carefully several times before beginning to work and especially before asking questions.
- Write your answers clearly and completely inside the box.

Question 1 (10 points). In class, we studied the pull algorithm (Cristian's algorithm) for physical clock synchronization in an asynchronous system, where each communication link has a minimum delay \min . Suppose the communication link from p to S has minimum delay \min_{ps} , and the communication link from S to p has minimum delay \min_{sp} . Consider a pull algorithm in this system, where the process p sends a request to S for the current time at S , and p can measure the round trip time T_{RT} until the response from S is received.

1. What should p set its clock to when it receives the response? Show your work.
2. What is the accuracy of p 's clock after it sets it? Show your work.



The earliest point p could receive the time t is $t + \min_{sp}$

The latest point is $t + T_{RT} - \min_{ps}$

The time interval is $[t + \min_{sp}, t + T_{RT} - \min_{ps}]$

$$t + T_{RT} - \min_{ps} - t - \min_{sp} = T_{RT} - \min_{ps} - \min_{sp}$$

① As a result, p should set its clock to

$$t + \min_{sp} + \frac{T_{RT} - \min_{ps} - \min_{sp}}{2} = t + \frac{T_{RT} + \min_{sp} - \min_{ps}}{2}$$

② The accuracy is $\pm \left(\frac{T_{RT} - \min_{ps} - \min_{sp}}{2} \right)$

Question 2 (10 points). Consider the matrix clocks in the Wu-Bernstein algorithm for the Replicated Log Problem. For each statement below, indicate whether the statement is TRUE or FALSE. If the statement is TRUE, provide a justification. If the statement is FALSE, describe a counter-example.

1. For any process p_i , it always holds that $T_i(i, k) \geq T_i(j, k)$ for $j = 1 \dots N$, $k = 1 \dots N$.
2. For any pair of processes p_i and p_j , $i \neq j$, it always holds that $T_i(i, i) \geq T_j(i, i)$.

① ~~The following~~ **Yes** The i th row of T_i refers to the direct knowledge of site i about ~~how~~ events in other sites, while other rows refer to the indirect knowledge.

For $j=i$, we have $T_i(i, k) = T_i(j, k)$ for all k .

For $j \neq i$, $T_i(i, k)$ will be updated ~~on~~ each receive event. We have

$$T_i(i, k) = \max[T_i(i, k), T_j(j, k)]$$

The indirect knowledge ~~is~~ is updated with other indirect knowledge.

$$T_i(r, s) = \max[T_i(r, s), T_j(r, s)]$$


Therefore, $T_i(i, k) \geq T_j(j, k)$ for $j, k = 1 \dots N$

② **Yes** $T_i(i, i)$ is recorded by i locally, so it always reflects the latest event ~~at~~ occurring at i , while $T_j(i, i)$ is an indirect knowledge about site i at site j . If site i sends a message to j everytime an update is made, then $T_i(i, i) = T_j(i, i)$. If there are many local events at i while i never tells j or any other site, $T_i(i, i) > T_j(i, i)$.


Therefore, $T_i(i, i) > T_j(i, i)$

Question 3 (10 points). Recall that in Lamport's mutual exclusion algorithm, a process accesses the resource when it has received GRANT from every ^{other} process and its request is at the head of its own priority queue. Suppose p_i has just started accessing the resource. Is it necessarily the case that p_i 's request is at the head of the priority queue at every process? Answer YES or NO and justify your answer.

NO. Think of the following counter-example.



In this case P_2 has just started accessing ~~the~~ ^{the} resource, but the priority queue of P_1 is $\boxed{P_1 P_2}$



P_2 will start as long as it receives grant. However, it's possible that a request is on the way when grant arrives. In that case, P_1 has already put itself at the head of queue and P_1 is waiting for a grant.